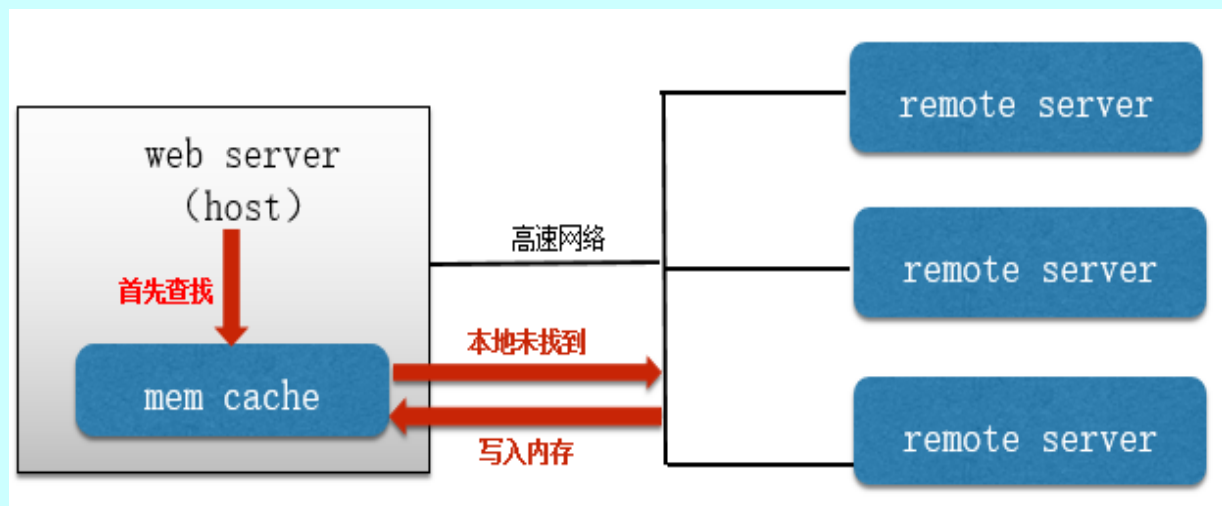


Cuckoo hash

软件学院《数据结构》讲义
内部使用



引例：Memcached 内存缓存



数据库中的商品信息

id (主键)	颜色	尺寸	价格
---------	----	----	----

将一件商品以键值对(key, value)的形式存储在内存里的cuckoo hash中，key和数据库中的主键id一致，能够唯一表示一件商品，value集合了颜色、尺寸、样式这些基本属性。

Cuckoo hash基本思想

Cuckoo hash的基本组成是2个hash 函数和一个hash table, 并且两个hash函数会确保将某个键映射至table中的不同位置, 也就是说对于任意键 k , $h1(k) \neq h2(k)$ 。一个键仅可能出现在table 中的 $h1(k)$ 位置或 $h2(k)$ 位置, 这两个位置中的唯一一个。

对比: 线性hash, 顺序遍历探测序列

链式hash, 需要遍历一次链表



Cuckoo

cuckoo意为布谷鸟, 布谷鸟会偷偷的在其它鸟的巢穴中产蛋, 当布谷鸟幼崽孵化出来后, 这些幼崽便会将其它幼鸟踢出巢穴, 以获得更大的生存空间。

Cuckoo hash基本操作



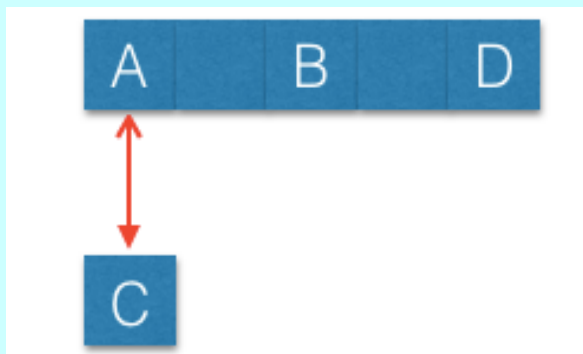
查找(get)操作：在cuckoo hash中，因为一个键仅可能出现在table中的 $h_1(k)$ 位置或者 $h_2(k)$ 位置，所以查找时仅需要探测这两个位置。

插入(put)操作：和其它hash 方法一样，cuckoo hash避免不了插入时的冲突。
对于某个键 k ，如果 $h_1(k)$ 位置发生冲突，则查看 $h_2(k)$ 位置，为空则将 k 插入至 $h_2(k)$ 位置，但如果 $h_2(k)$ 非空呢？

Cuckoo hash基本操作

插入(put)操作：和其它hash 方法一样，cuckoo hash避免不了插入时的冲突。

对于某个键k，如果 $h1(k)$ 位置发生冲突，则查看 $h2(k)$ 位置，
为空则将k插入至 $h2(k)$ 位置，但如果 $h2(k)$ 非空呢？



问题：最后一个被踢出的元素永远无法找到一个空位置，这样整个踢出过程便无法终止。无法终止的踢出过程都会形成一个环。

键值分离存储

hash table中每一项存储的是
<key, **address**>, address记录
了值所在的地址。

```
typedef int KeyType;
class Cuckoo{
protected:
    std::mutex mtx;
    KeyType T[SIZE];
    // hash key by hash func 1
    int hash1(const KeyType &key);
    // hash key by hash func 2
    int hash2(const KeyType &key);
    // find key by hash func 1 in T, exist return key otherwise 0
    KeyType get1(const KeyType &key);
    // find key by hash func 2 in T, exist return key otherwise 0
    KeyType get2(const KeyType &key);
    void bt_evict(const KeyType &key, int which, int pre_pos);
public:
    Cuckoo();
    ~Cuckoo();
    KeyType get(const KeyType &key);
    void put(const KeyType &key);
};
```

串行Get

```
Cuckoo::Cuckoo() {
    memset(T, 0, sizeof(KeyType) * SIZE);
}

//~Cuckoo();

int Cuckoo::hash1(const KeyType &key) {
    assert(SIZE != 0);
    int half_siz = SIZE / 2;
    return key%half_siz;
}

int Cuckoo::hash2(const KeyType &key) {
    assert(SIZE != 0);
    int half_siz = SIZE / 2;
    return key/half_siz%half_siz + half_siz;
}
```

```
// find key by hash func 1 in T, exist return key otherwise 0
KeyType Cuckoo::get1(const KeyType &key) {
    return (T[hash1(key)] == key)?key:0;
}

// find key by hash func 2 in T, exist return key otherwise 0
KeyType Cuckoo::get2(const KeyType &key) {
    return (T[hash2(key)] == key)?key:0;
}

KeyType Cuckoo::get(const KeyType &key) {
    // 0 is reserved for null, invalid input
    if(key == 0) {
        printf("invalid key\n");
        return 0;
    }

    KeyType result = get1(key);
    if(result == 0) {
        result = get2(key);
    }

    return result;
}
```

并行Get

共享变量为何没有锁
保护?



```
static const int TOTAL = 10;

int main(int argc, char* argv[]) {
    Cuckoo test;
    // single-thread to put [1, TOTAL]
    for(int i = 1; i <= TOTAL; ++i) {
        test.put(i);
    }

    // create multiple threads to get in parallel
    std::vector<std::thread> threads;
    threads.clear();
    for(int i = 1; i <= TOTAL; ++i) {
        threads.emplace_back([&](int thread_id) {
            printf("thread: %d get %d\n", thread_id, test.get(thread_id));
        }, i);
    }
    for(int i = 0; i < TOTAL; ++i) {
        threads[i].join();
    }
    return 0;
}
```


Put

```
template <typename T>
inline void swap(T* a, T* b) {
    assert(a != NULL && b != NULL);
    T tmp = *a;
    *a = *b;
    *b = tmp;
}

void Cuckoo::put(const KeyType &key) {
    if(key == 0) {
        printf("invalid key\n");
        return;
    }
    if(get(key) != 0) {
        printf("duplicate key, put fail\n");
        return;
    }
    // basic way
    if(T[hash1(key)] == 0) {
        T[hash1(key)] = key;
    } else if(T[hash2(key)] == 0) {
        T[hash2(key)] = key;
```

```
} else { // two place for one certain key has been occupied, need evict others
    // basic way
    KeyType evicted = key;
    // determine which pos hash1 or hash2 to put key
    // 0 is hash1, 1 is hash2
    int which = 0;
    // first evict key in hash1
    int idx = hash1(evicted);
    // != 0 means place has been occupied
    // if there is a cycle, maybe cannot terminate
    int pre_pos = -1;
    while(T[idx] != 0) {
        printf("evicted key %d from %d to %d\n", evicted, pre_pos, idx);
        swap(&evicted, &T[idx]);
        pre_pos = idx;
        which = 1 - which;
        idx = (which == 0)?hash1(evicted):hash2(evicted);
    }
    printf("evicted key %d from %d to %d\n", evicted, pre_pos, idx);
    T[idx] = evicted;
}
```

基于回溯的实现

当 $n \neq 1$ 时，程序会不断的向下调用，形成一个没有分叉的递归调用树。

当 $n = 1$ 时，程序从调用树的叶子节点返回计算结果，并且每一层都会向调用层返回自己这一层的计算结果，到达根节点时便会得到最终结果

```
int fac(int n) {  
    if(n==1)  
        return n;  
    else  
        return n * fac(n-1);  
}
```

基于回溯的实现

假定产生的踢出序列为A->B->C->D->nil

先踢出D，依次向上直到A，我们便可以发现：在保证将某个键插入到指定位置的操作是原子的前提下，就可以确保这些元素始终在hash table里

```
void Cuckoo::bt_evict(const KeyType &key, int which, int pre_pos) {
    int idx = (which == 0)?hash1(key):hash2(key);
    // basic case: find a empty pos for the last evicted element
    if(T[idx] == 0) {
        printf("evicted key %d from %d to %d\n", key, pre_pos, idx);
        T[idx] = key;
        return;
    }
    printf("evicted key %d from %d to %d\n", key, pre_pos, idx);
    KeyType cur = T[idx];
    // first evict latter elements
    bt_evict(cur, 1 - which, idx);
    T[idx] = key;
}
```

基于回溯的实现

```
void Cuckoo::put(const KeyType &key) {
    if(key == 0) {
        printf("invalid key\n");
        return;
    }
    if(get(key) != 0) {
        printf("duplicate key, put fail\n");
        return;
    }
    // basic way
    if(T[hash1(key)] == 0) {
        T[hash1(key)] = key;
    } else if(T[hash2(key)] == 0) {
        T[hash2(key)] = key;
    } else { // two place for one certain key has been occupied, need evict others
        // backtrace way
        bt_evict(key, 0, -1);
    }
}
```

并行Put

```
void Cuckoo::put(const KeyType &key) {
    if(key == 0) {
        printf("invalid key\n");
        return;
    }
    if(get(key) != 0) {
        printf("duplicate key, put fail\n");
        return;
    }
    // basic way
    if(T[hash1(key)] == 0) {
        T[hash1(key)] = key;
    } else if (T[hash2(key)] == 0) {
        T[hash2(key)] = key;
    }
```

```
} else { // two place for one certain key has been occupied, need evict others
    // lock way
    // need lock for write-operations
    std::unique_lock<std::mutex> lck(mtx);

    KeyType evicted = key;
    // determine which pos hash1 or hash2 to put key
    // 0 is hash1, 1 is hash2
    int which = 0;
    // first evict key in hash1
    int idx = hash1(evicted);
    // != 0 means place has been occupied
    // if there is a cycle, maybe cannot terminate
    int pre_pos = -1;
    while(T[idx] != 0) {
        printf("evicted key %d from %d to %d\n", evicted, pre_pos, idx);
        swap(&evicted, &T[idx]);
        pre_pos = idx;
        which = 1 - which;
        idx = (which == 0)?hash1(evicted):hash2(evicted);
    }
    printf("evicted key %d from %d to %d\n", evicted, pre_pos, idx);
    T[idx] = evicted;
}
}
```

死循环

```
#include <vector>

#include "cuckoo.cpp"

using namespace cuckoo;

static const int TOTAL = 28;

int main(int argc, char* argv[]) {
    Cuckoo test;



    // single-thread to put [1, TOTAL]
    for(int i = 1; i <= TOTAL; ++i) {
        test.put(i);
    }

    return 0;
}
```


```

1 evicted key 28 from -1 to 3
2 evicted key 3 from 3 to 25
3 evicted key 2 from 25 to 2
4 evicted key 27 from 2 to 26
5 evicted key 26 from 26 to 1
6 evicted key 1 from 1 to 25
7 evicted key 3 from 25 to 3
8 evicted key 28 from 3 to 26
9 evicted key 27 from 26 to 2
10 evicted key 2 from 2 to 25
11 evicted key 1 from 25 to 1
12 evicted key 26 from 1 to 26
13 evicted key 28 from 26 to 3
14 evicted key 3 from 3 to 25
15 evicted key 2 from 25 to 2

```



可以终止的环



只有当一个键的两个可选位置都各自形成一个环结构时，才会导致整个过程无法终止

检测循环路径的方法也比较简单，可以预先设定一个阈值(threshold)，当循环次数或者递归调用次数超过阈值时，就可以认为产生了循环路径。一旦发生循环路径之后，常规方法就是进行rehash操作

New hash functions are chosen, and the whole data structure is **rebuilt** (“rehashed”)

性能分析

Cuckoo hash的总容量限制为500个键

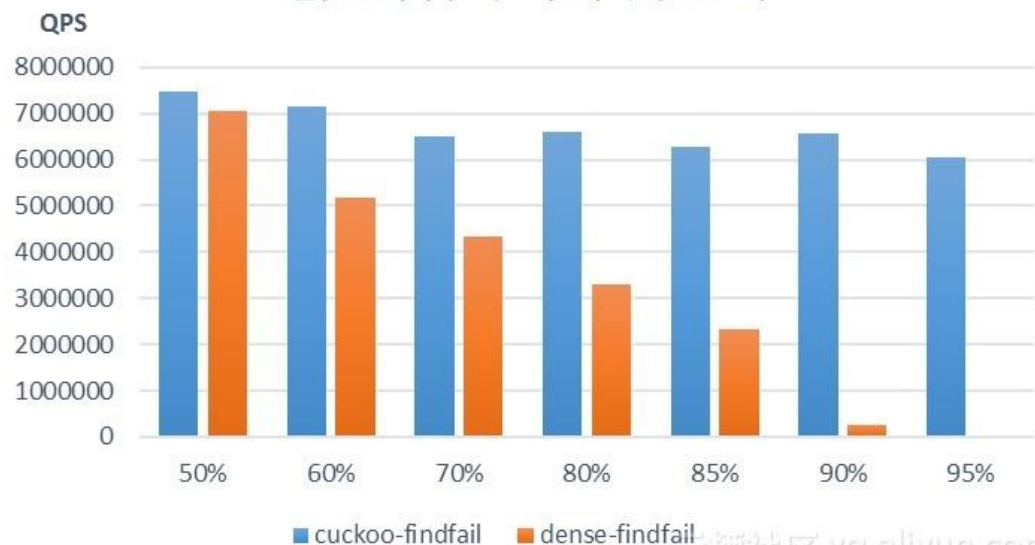
查找键总数 Hash 方法	50	250	375	500
Cuckoo hash	1	1	1.33	1.5
链式hash	1	3	4.67	5.5

最多只会访问两个位置的键，所以每次的比较次数不会超过2，平均比较次数当然就在2以内。从表中也可以发现，当cuckoo hash的负载因子分别为0.10、0.50、0.75、1.00时，平均比较次数都维持在2以内


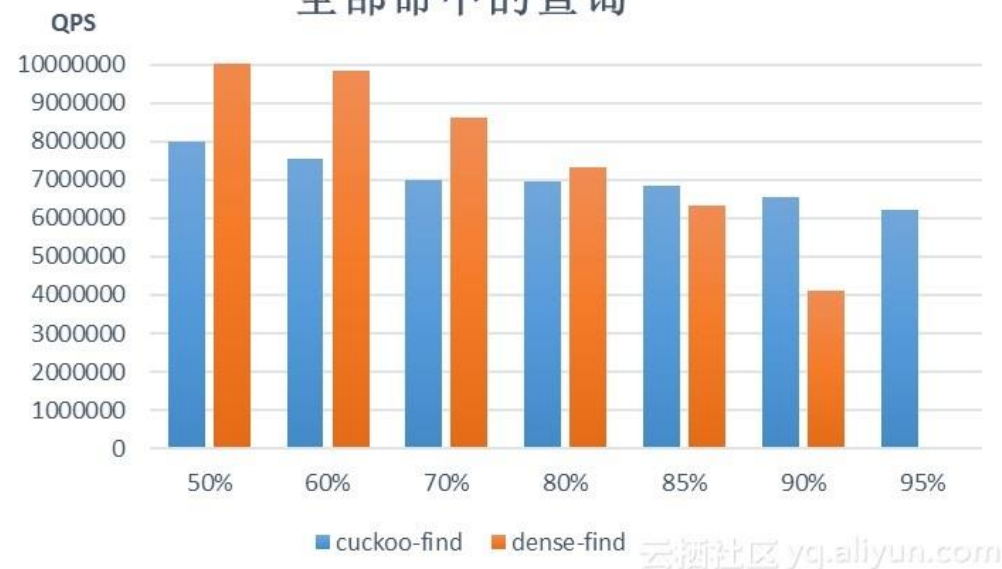
$$\text{负载因子} = \frac{\text{键的总数}}{\text{总容量}}$$

性能分析

绝大部分不命中的查询

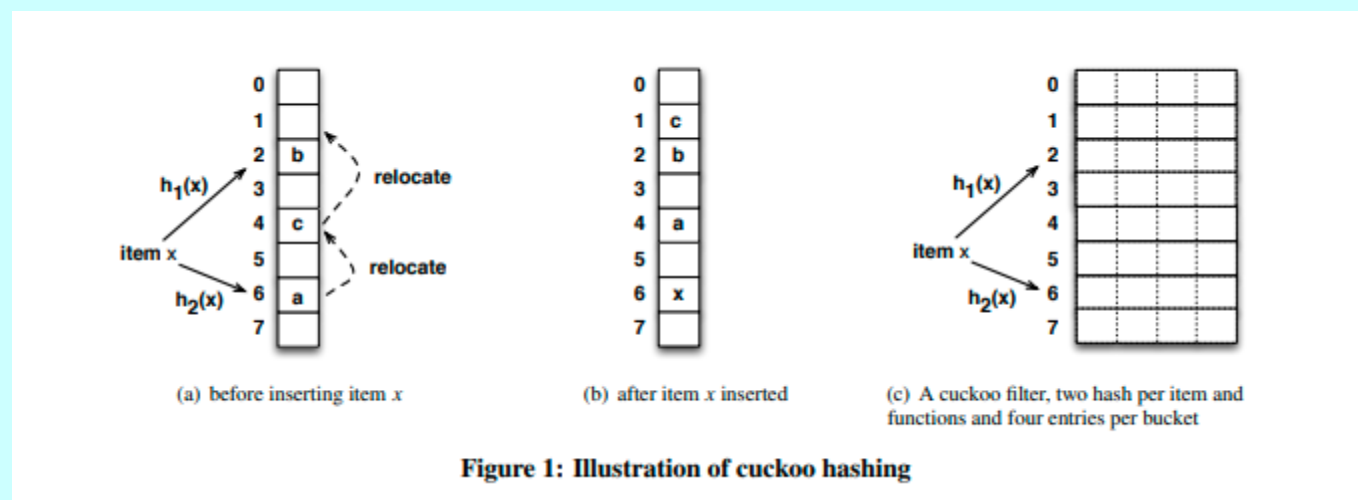


全部命中的查询



Source: <https://developer.aliyun.com/article/563053>

扩展: Cuckoo Filter



Cuckoo Filter: Practically Better Than Bloom

Bin Fan, David G. Andersen, Michael Kaminsky[†], Michael D. Mitzenmacher[‡]
Carnegie Mellon University, [†]Intel Labs, [‡]Harvard University
{binfan,dga}@cs.cmu.edu, michael.e.kaminsky@intel.com, michaelm@eecs.harvard.edu

Next

- 并行最小生成树
- 数据结构讲义